

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
In [2]: df = pd.read_csv("Iris.csv")
# .dropna(axis="columns")
df
```

```
Out[2]:
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa
...
145	146	6.7	3.0	5.2	2.3	Iris-virginica
146	147	6.3	2.5	5.0	1.9	Iris-virginica
147	148	6.5	3.0	5.2	2.0	Iris-virginica
148	149	6.2	3.4	5.4	2.3	Iris-virginica
149	150	5.9	3.0	5.1	1.8	Iris-virginica

150 rows × 6 columns

```
In [3]: df.head()
```

```
Out[3]:
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa

Data cleaning and pre processing

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                    150 non-null   int64
1   SepalLengthCm         150 non-null   float64
2   SepalWidthCm          150 non-null   float64
3   PetalLengthCm         150 non-null   float64
4   PetalWidthCm          150 non-null   float64
5   Species               150 non-null   object
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
```

```
In [5]: df.describe()
```

```
Out[5]:
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
count	150.000000	150.000000	150.000000	150.000000	150.000000
mean	75.500000	5.843333	3.054000	3.758667	1.198667
std	43.445368	0.828066	0.433594	1.764420	0.763161
min	1.000000	4.300000	2.000000	1.000000	0.100000
25%	38.250000	5.100000	2.800000	1.600000	0.300000
50%	75.500000	5.800000	3.000000	4.350000	1.300000
75%	112.750000	6.400000	3.300000	5.100000	1.800000
max	150.000000	7.900000	4.400000	6.900000	2.500000


```
In [6]: df.columns
```

```
Out[6]: Index(['Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm',
              'Species'],
              dtype='object')
```

EDA and VISUALIZATION

```
In [7]: sns.pairplot(df)
```


```
Out[7]: <seaborn.axisgrid.PairGrid at 0x25a60c28790>
```



```
In [8]: sns.distplot(df[ 'SepalLengthCm' ])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)


```
Out[8]: <AxesSubplot:xlabel='SepalLengthCm', ylabel='Density'>
```



```
In [9]: df1 = df[['Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm',  
              'Species']]
```

```
In [10]: sns.heatmap(df1.corr())
```

Out[10]: <AxesSubplot:>



```
In [11]: x = df1[['Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalWidthCm']]  
         y = df1['PetalLengthCm']
```

split the data into training and test data

```
In [12]: x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.3)
```

```
In [13]: lr = LinearRegression()  
         lr.fit(x_train, y_train)
```

Out[13]: LinearRegression()

In [14]: `lr.intercept_`

Out[14]: -0.3308682501944493


In [15]: `coeff = pd.DataFrame(lr.coef_, x.columns, columns = ['Co-efficient'])`
`coeff`

Out[15]:

	Co-efficient
Id	0.002636
SepalLengthCm	0.750799
SepalWidthCm	-0.688092
PetalWidthCm	1.328811

In [16]: `prediction = lr.predict(x_test)`
`plt.scatter(y_test, prediction)`

Out[16]: <matplotlib.collections.PathCollection at 0x25a624598b0>



In [17]: `lr.score(x_test,y_test)`

Out[17]: 0.9571365220360376

ACURACY

In [18]: `from sklearn.linear_model import Ridge,Lasso`

In [19]: `rr=Ridge(alpha=10)`
`rr.fit(x_train,y_train)`

```
rr.score(x_test,y_test)  
rr.score(x_train,y_train)
```

Out[19]: 0.9544588048793379

In [20]:

```
rr.score(x_test,y_test)
```

Out[20]: 0.9395084922916778

In [21]:

```
la = Lasso(alpha=10)  
la.fit(x_train,y_train)
```

Out[21]: Lasso(alpha=10)

In [22]:

```
la.score(x_test,y_test)
```

Out[22]: 0.7302717056543535