

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
In [2]: df = pd.read_csv("Sleep.csv")
# .dropna(axis="columns")
df
```

Out[2]:

	Person ID	Gender	Age	Occupation	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	BMI Category	Blood Pressure	Heart Rate
0	1	Male	27	Software Engineer	6.1	6	42	6	Overweight	126/83	
1	2	Male	28	Doctor	6.2	6	60	8	Normal	125/80	
2	3	Male	28	Doctor	6.2	6	60	8	Normal	125/80	
3	4	Male	28	Sales Representative	5.9	4	30	8	Obese	140/90	
4	5	Male	28	Sales Representative	5.9	4	30	8	Obese	140/90	
...
369	370	Female	59	Nurse	8.1	9	75	3	Overweight	140/95	
370	371	Female	59	Nurse	8.0	9	75	3	Overweight	140/95	
371	372	Female	59	Nurse	8.1	9	75	3	Overweight	140/95	
372	373	Female	59	Nurse	8.1	9	75	3	Overweight	140/95	
373	374	Female	59	Nurse	8.1	9	75	3	Overweight	140/95	

374 rows × 13 columns



```
In [3]: df.head()
```

Out[3]:

	Person ID	Gender	Age	Occupation	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	BMI Category	Blood Pressure	Heart Rate
0	1	Male	27	Software Engineer	6.1	6	42	6	Overweight	126/83	77

	Person ID	Gender	Age	Occupation	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	BMI Category	Blood Pressure	Heart Rate
1	2	Male	28	Doctor	6.2	6	60	8	Normal	125/80	75
2	3	Male	28	Doctor	6.2	6	60	8	Normal	125/80	75
3	4	Male	28	Sales Representative	5.9	4	30	8	Obese	140/90	85
4	5	Male	28	Sales Representative	5.9	4	30	8	Obese	140/90	85

Data cleaning and pre processing

In [4]:

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 374 entries, 0 to 373
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Person ID                            374 non-null    int64
1   Gender                               374 non-null    object
2   Age                                   374 non-null    int64
3   Occupation                           374 non-null    object
4   Sleep Duration                       374 non-null    float64
5   Quality of Sleep                     374 non-null    int64
6   Physical Activity Level              374 non-null    int64
7   Stress Level                         374 non-null    int64
8   BMI Category                         374 non-null    object
9   Blood Pressure                       374 non-null    object
10  Heart Rate                           374 non-null    int64
11  Daily Steps                          374 non-null    int64
12  Sleep Disorder                       374 non-null    object
dtypes: float64(1), int64(7), object(5)
memory usage: 38.1+ KB
```

In [5]:

```
df.describe()
```

Out[5]:

	Person ID	Age	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	Heart Rate	Daily Steps
count	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000
mean	187.500000	42.184492	7.132086	7.312834	59.171123	5.385027	70.165775	6816.844920
std	108.108742	8.673133	0.795657	1.196956	20.830804	1.774526	4.135676	1617.915679
min	1.000000	27.000000	5.800000	4.000000	30.000000	3.000000	65.000000	3000.000000
25%	94.250000	35.250000	6.400000	6.000000	45.000000	4.000000	68.000000	5600.000000
50%	187.500000	43.000000	7.200000	7.000000	60.000000	5.000000	70.000000	7000.000000
75%	280.750000	50.000000	7.800000	8.000000	75.000000	7.000000	72.000000	8000.000000
max	374.000000	59.000000	8.500000	9.000000	90.000000	8.000000	86.000000	10000.000000

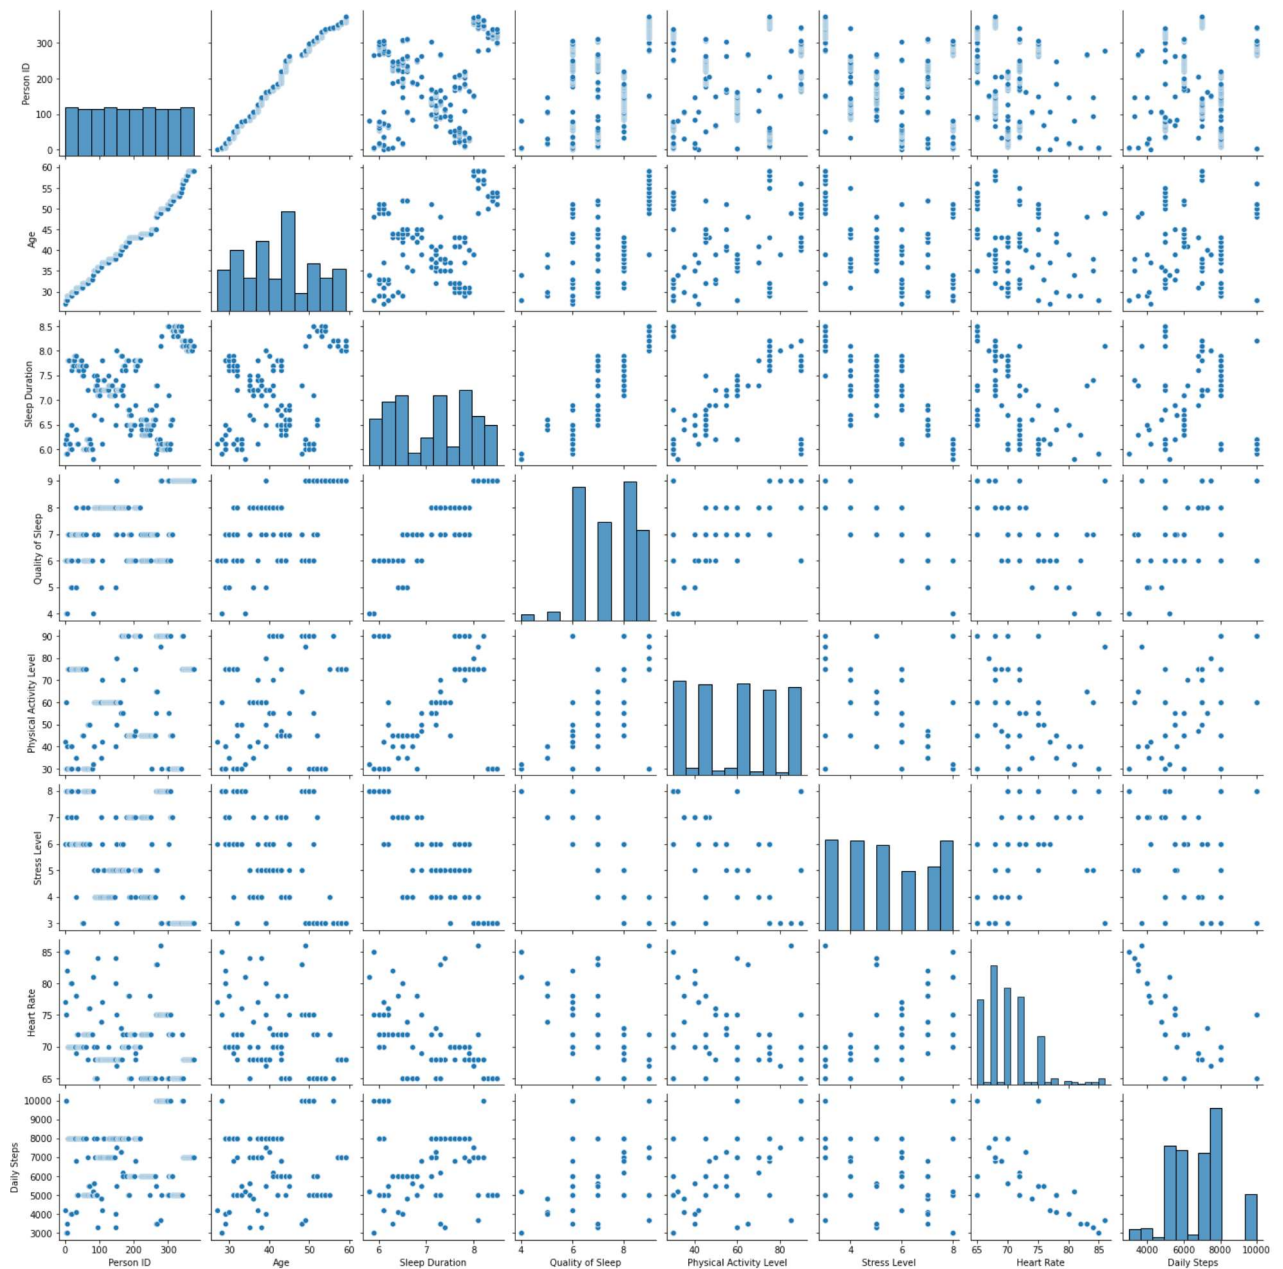
```
In [6]: df.columns
```

```
Out[6]: Index(['Person ID', 'Gender', 'Age', 'Occupation', 'Sleep Duration',
              'Quality of Sleep', 'Physical Activity Level', 'Stress Level',
              'BMI Category', 'Blood Pressure', 'Heart Rate', 'Daily Steps',
              'Sleep Disorder'],
              dtype='object')
```

EDA and VISUALIZATION

```
In [7]: sns.pairplot(df)
```


```
Out[7]: <seaborn.axisgrid.PairGrid at 0x1f053131ca0>
```



```
In [8]: sns.distplot(df['Sleep Duration'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)


```
Out[8]: <AxesSubplot:xlabel='Sleep Duration', ylabel='Density'>
```



```
In [9]: df1 = df[['Person ID', 'Gender', 'Age', 'Occupation', 'Sleep Duration',  
                'Quality of Sleep', 'Physical Activity Level', 'Stress Level',  
                'BMI Category', 'Blood Pressure', 'Heart Rate', 'Daily Steps',  
                'Sleep Disorder']]
```

```
In [10]: sns.heatmap(df1.corr())
```

```
Out[10]: <AxesSubplot:>
```



```
In [11]: x = df1[['Person ID', 'Sleep Duration',
               'Quality of Sleep', 'Stress Level',
               'Heart Rate', 'Daily Steps']]
         y = df1['Age']
```

split the data into training and test data

```
In [12]: x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.3)
```

```
In [13]: lr = LinearRegression()
         lr.fit(x_train, y_train)
```

```
Out[13]: LinearRegression()
```

```
In [14]: lr.intercept_
```

```
Out[14]: 11.899286331704307
```

```
In [15]: coeff = pd.DataFrame(lr.coef_, x.columns, columns=['Co-efficient'])
         coeff
```


```
Out[15]:
```

	Co-efficient
Person ID	0.077905
Sleep Duration	0.287208
Quality of Sleep	0.591638
Stress Level	0.031027

Co-efficient	
Heart Rate	0.123934
Daily Steps	0.000072

```
In [16]: prediction = lr.predict(x_test)
plt.scatter(y_test, prediction)
```

```
Out[16]: <matplotlib.collections.PathCollection at 0x1f05513ed00>
```



```
In [17]: lr.score(x_test,y_test)
```

```
Out[17]: 0.9819159155441144
```

ACURACY

```
In [18]: from sklearn.linear_model import Ridge,Lasso
```

```
In [19]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
rr.score(x_test,y_test)
rr.score(x_train,y_train)
```

```
Out[19]: 0.9868875766465539
```

```
In [20]: rr.score(x_test,y_test)
```

```
Out[20]: 0.9819248705055501
```

```
In [21]: la = Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[21]: Lasso(alpha=10)

In [22]: `la.score(x_test,y_test)`

Out[22]: 0.9779820947092206